Narayana Rongali

Data Engineer
rongalisssnarayana@gmail.com | +91-6300094534 | linkedin.com/in/narayanarongali

Summary

Data Engineer with 3+ years designing batch and incremental data platforms on Databricks, Azure and AWS for enterprise retail and CPG clients (PepsiCo FLNA, Adobe). Owns pipelines end to end, from ingestion and dimensional modelling to orchestration, data quality and BI serving, processing 1.2M+ records daily across 5 retailers. Strong record in performance and cost: Spark runtime cut ~80% and cluster cost ~35%. Databricks Certified Data Engineer Professional.

Technical skills

- **Languages:** SQL (T-SQL), Python, PySpark, pandas
- **Processing and orchestration:** Databricks, Delta Lake, Auto Loader, Azure Data Factory, Airflow, AWS Glue, Kafka and Debezium (CDC), Apache Iceberg
- **Data stores:** Snowflake, SQL Server, Azure SQL, Synapse, Redshift, Couchbase (NoSQL), Teradata, IBM DB2, MySQL
- **Cloud and DevOps:** AWS (S3, EC2, IAM), Azure (ADLS Gen2, Key Vault, Monitor, AKS), Azure DevOps CI/CD, ARM templates, Git
- **Modelling:** Dimensional modelling (star schema, conformed dimensions, SCD Type 2), medallion architecture, incremental/CDC loads
- **BI and quality:** Power BI, DAX, data quality frameworks, reconciliation, root cause analysis, Spark and Snowflake performance tuning

Professional experience

Senior Analyst, Data Engineering | LatentView Analytics, Chennai (Sep 2023 - Present)

Client: PepsiCo (Frito-Lay North America), CPO order generation and forecasting

- Built an automated order-generation and forecasting platform for 5 retailers (Sam's Club, Walmart, BJ's, Target, Amazon) on ADF and Databricks, processing 1.2M+ store-item records daily; order cycle cut from ~4 hrs to under 20 min.
- Designed a layered architecture: ADLS Gen2/Delta medallion for batch, a star-schema serving layer for Power BI, and Couchbase as the operational store for real-time order write-backs from a field app on Azure AKS, isolating analytical and operational workloads.
- Implemented incremental, CDC-style loads (watermarks, Delta MERGE upserts, Auto Loader schema evolution) and reconciled Couchbase write-backs to SQL Server: duplicate and malformed records down 85%.
- Modelled a star schema at store-item-day grain with conformed dimensions (store, item, retailer, date) and SCD Type 2 history: Power BI query time improved 55%; 10+ dashboards with automated refresh cut manual reporting 70%.
- Orchestrated dependency-aware Airflow DAGs (retries, SLAs, backfills) across ingestion, Databricks transforms and data quality gates, so a failed check blocks downstream order generation.
- Cut the nightly Spark run from ~3.5 hrs to 40 min (~80%) and cluster cost ~35% by fixing skew (salting, broadcast joins, AQE) and file layout (partitioning, OPTIMIZE, Z-ORDER).
- Optimised Snowflake queries and warehouse usage (query profile analysis, partition pruning, clustering keys, warehouse sizing, result caching) to lower runtime and credit consumption.
- Built a PySpark data quality framework (nulls, ranges, referential integrity, reconciliation) over 25+ tables with Azure Monitor alerting: downstream incidents down 45%.
- Migrated the platform and on-premises SQL Server 2019 workloads to AWS (S3, EC2, IAM), rebuilding stored procedures and views; key forecast procedure cut from ~45 to 6 min.
- Built metadata-driven ADF pipelines (incl. SHIR-based API ingestion): retailer onboarding from ~2 weeks to 3 days. Automated releases via Azure DevOps and ARM templates with Key Vault secrets: release time down 50%.

Client: Adobe, data engineering support

- Supported ETL pipelines loading business data into reporting tables; wrote and tuned the SQL transformations.
- Ran data quality checks (row counts, nulls, duplicates), resolved pipeline failures and documented pipeline logic and data definitions.

Certifications and awards

- Databricks Certified Data Engineer Professional (2026); Databricks Certified Generative AI Engineer Associate (2026)
- Spot Award (Going Above and Beyond) and Encore Award (Excellence in Customer Centricity), LatentView Analytics

Education

B.Tech in Chemical Engineering, Indian Institute of Petroleum and Energy, Visakhapatnam (Aug 2019 - Apr 2023)